University of Warsaw
Faculty of Mathematics, Informatics and Mechanics

Marta Nowakowska

Student no. 385914

Time series prediction using self-attention models

Master's thesis
in COMPUTER SCIENCE

Supervisor:
dr hab. Marek Cygan, prof. UW
University of Warsaw

Co-supervisor:
dr hab. Piotr Miłoś
Institute of Mathematics, Polish Academy of Sciences

Warsaw, September 2022

Abstract

The Transformer model has excelled at various sequence-processing problems. However, the past works dedicated to time series forecasting with neural networks commonly simplified this task to predicting parameters of simple distributions, i.e. the Gaussian distribution. This thesis explores the idea of applying the Transformer to forecasting parameters of a granular categorical distribution for each predicted data point. This model is compared against the Gaussian distribution-based Transformer in the presented experiments.

Keywords

Time series, Neural network, Transformer, Self-attention model

Thesis domain (Socrates-Erasmus subject area codes)

11.3 Informatics, Computer Science

Subject classification

I. Computing Methodologies
I.2. Artificial Intelligence
I.2.6. Learning
Connectionism and neural nets

Tytuł pracy w języku polskim

Prognozowanie szeregów czasowych przy użyciu modeli z uwagą

Contents

Introduction

Time series and their importance

Forecasting future values and their probability distribution on the basis of historical datapoints is one of the most prevalent problems in both industry and research. For example, time series analysis can be used for predicting the number of ride-hailing trips completed during special events like holidays ([10]), but also for estimating stock market behavior ([11]) and solving resource allocation tasks ([12]).

This wide usage brought time series a lot of attention in the past. One of the most known methods for their forecasting are based on Box-Jenkins method like the ARIMA (autoregressive integrated moving average) [8] model, but there also exist ones using the exponential smoothing techniques and state space models [9]. Despite their popularity, they come with some fundamental disadvantages. For example, they need to be fit separately for every similar sequence ([1], [7]) and cannot be trained to recognize these similarities which makes using them time-inefficient in the long term. Furthermore, these models typically use parameters that have to be handpicked by the user ([1], [7]), which makes them more suitable for field experts with knowledge of both parameter selection procedure and preferably time series' behavior.

The aforementioned issues are not a problem for deep neural networks, since they can be trained to recognize a group of similar time series and their hyper-parameters have to be only selected once for the training and not at the prediction time. What is more, once trained neural network can be adapted to new data using transfer learning.

On top of that, neural networks proved to be very successful in other fields, including natural language translation, which is a problem with similar properties to the time series forecasting. In particular, the translation can be defined as a mapping of one sequence to a second one, similarly to how one can map historical values to the future ones. Thus, the achievements of natural language processing (NLP) field, including RNN models have almost directly transferred to the time series forecasting. However, RNNs and their successors (LSTM [2], GRU [3]) have difficulty with long sequences [4], causing them to struggle with time series that can contain long-term dependencies.

The next important step in the development of deep neural network adaptations for the time series forecasting was the Transformer [5]. Instead of using recurrent networks it relies almost entirely on the self-attention mechanism, which can be trained to estimate dependencies in arbitrarily long sequences, making it possible for the model to deal with larger-scale forecasting. However, there are several problems with integrating it directly for the time series forecasting usage. One of them is that the scaled dot-product attention can only capture dependencies between particular data points, but it cannot learn the time series' shape, which makes it struggle with capturing its seasonality and other patterns or anomalies ([1]). The second important issue is that its space complexity grows quadratically with the sequence length, making it difficult to train larger Transformer models on long time series ([1]). One

solution to these problems was proposed in [1] by introducing convolutional self-attention and LogSparse Transformer.

Problem statement

Previously neural networks including the Transformer were commonly used for predicting parameters of a simple distribution like Gaussian per each forecasted point in time ([7], [1]). This work expands that idea to modeling any continuous distribution by approximating it with a categorical one. We present a comparison of both methods on commonly used datasets and metrics.

Contributions

Work described in this thesis has been done via a collaboration with scientists from Institute of Mathematics of Polish Academy of Sciences.

My most important contributions are:

1. Gaussian distribution-based Transformer

2. Experiments with the Gaussian distribution-based Transformer

3. Base implementation of the dataset feeder

4. Time features as in [1] (on the data feeding side)

Chapter 1

Background

1.1. Time series and notation

We define time series as sequence of real numbers measured across time with a constant **frequency**. Frequency is defined in time units, e.g. weeks, days, hours. For multiple time series represented as a matrix $\mathbf{y}$, where each row is a separate time series, $y_t^{(i)}$ denotes t time step of a series number i.

One can extend the definition above to a sequence of any data type, but it is not necessary for the datasets that we use for this thesis.

1.2. Transformer model architecture

Transformer [5] is a deep neural network model that relies on a scaled dot-product attention mechanism introduced in the same paper. There exist some modifications to the model presented in [5], but we follow the "decoder-only" [6] architecture type.

1.2.1. Overview

Our implementation of the Transformer model consists mainly of a causal convolution layer, positional encoding layer, N identical decoder layers, layer normalization and a fully-connected layer. Most of the layers' outputs have d_{model} dimension. I present the general architecture on Figure 1.1.

Positional encoding allows us to feed the model embedded information about how the data points are ordered in the inputted time series. We use a trainable embeddings for the positional encodings and add it to the input before it reaches the first decoder layer. [pk: ???] [mn: I changed it a little, but I am not sure whether it is still confusing.]

A single **decoder** layer consists of a causal attention sub-layer and a feed-forward sub-layer:

- The feed-forward sub-layer contains a simple feed-forward neural network with two dense layers and a ReLU activation in-between.

- The causal attention sub-layer contains the attention mechanism that I describe later on.

Both sub-layers are residual and wrapped with a layer normalization and a dropout layer.
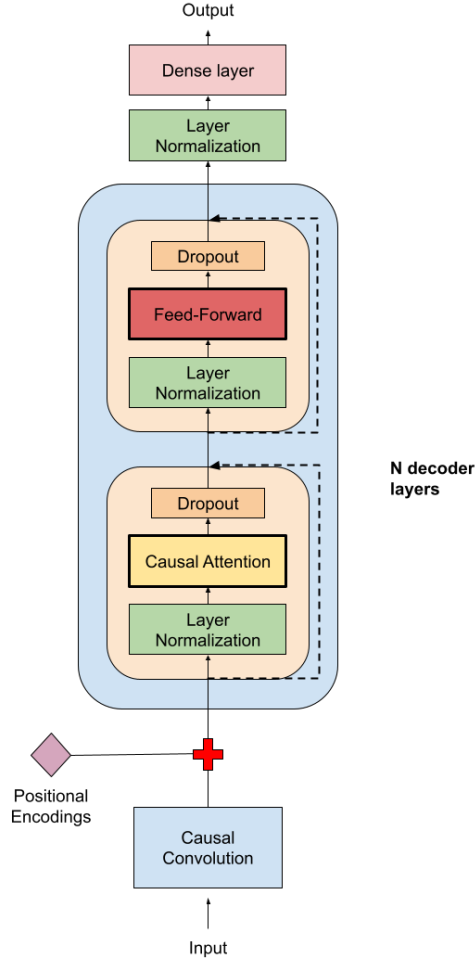
Figure 1.1: Transformer's architecture. The dotted lines represent residual connections.

1.3. Attention

1.3.1. Scaled Dot-Product Attention

Attention is the key part to explaining how the Transformer works. It enables us to train the neural network to encode the relations between different parts of the input (e.g., a time series' data points).

Scaled dot-product attention introduced along with the Transformer architecture in [5] is perhaps the most known kind of attention function. To calculate the attention score vector between one data point and all that come after, one first needs to provide a query vector for the first data point and a set of vector pairs consisting of keys and values for all the other data points. We match our query against all of the keys by a dot product and use the softmax function to normalize the result. Finally, by multiplying the softmax results with the corresponding values we get a set of weighted values vectors. After summing them we get the attention score vector for one data point.

One can informally see these calculations as operations on a neural network-approximated dictionary, where each key corresponds to a certain value and we treat the dot operation between the query and keys as a measure of similarity. When they are orthogonal, the dot product will be equal to 0. On the other hand, for vectors of fixed length the dot product

is maximal when they have the same direction. When the result of softmax is large between two data points, we can informally say that the attention relates these two points.

In reality we do not need to calculate the attention sequentially for every data point, but rather use matrix operations. One needs to provide three matrices to the attention layer: $Q, K \in \mathbb{R}^{d_I \times d_k}, V \in \mathbb{R}^{d_I \times d_v}$ (respectively called queries, keys and values). They are created by multiplying the input matrix $I \in \mathbb{R}^{d_I \times d_{model}}$ with three distinct trainable weight matrices $W_Q \in \mathbb{R}^{d_{model} \times d_k}, W_K \in \mathbb{R}^{d_{model} \times d_k}, W_V \in \mathbb{R}^{d_{model} \times d_v}$.

For given Q, K, V we compute the attention with the following formula:

$$Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_k}} \cdot M)V \text{ (from [1])}$$

[pk: a picture is worth a thousand words] [mn: I think here maybe it's worth adding some words, since I did not talk a lot about the informal understanding of attention. Above I added some paragraphs about how to calculate and understand it. If I were to add a picture, maybe the one from the paper (Figure 2 in Attention...)? But I think this formula is very short anyways.]

where d_k stands for the dimension of Q and K. Dividing by its square root is supposed to balance out the growth of Q and K^T dot product that occurs with larger values of d_k.

M is a **mask matrix** that sets all the elements corresponding to connections between future and past (upper triangular) to $-\infty$. It prevents the model from attempting to relate the past events to the future ones. [pk: show the matrix] [mn: I fixed it right now, since I see I was wrong before of how it should look like. Not sure whether you'd still want me to show it, since it should be one of the simplest matrices spoken about here. WDYT?]

While calculating the softmax score for any two data points in a time series, the scaled dot-product attention does not use any information about other data points. We call this property **context-blindness**. Sometimes this might lead the model to assigning strong attention score between events that have similar values, even if their contexts (neighboring data points) are very different. That might cause the model to incorrectly recognize patterns in the time series ([1]).

1.3.2. Convolutional Attention

Definition 1.3.1 *Causal convolution is a convolution used on input padded with $k - 1$ zeros at the start, where k denotes the size of the kernel.*

This type of convolution is particularly useful for time series, since it enables us to make predictions with only partial input, essentially generating a completely new time series.

The convolutional attention [1] is calculated in almost the same way as the scaled dot-product attention. It only differs in the way it computes queries and keys. Instead of applying a matrix multiplication with weight matrices, it uses a causal convolution with a stride of 1. When the kernel's size is equal to 1 it is equal to the matrix multiplication.

The usage of convolution is supposed to offset the dot-product attention's context-blindness. Since the convolution layer learns to recognize input's context, stronger attention score should be assigned to similar events.

1.3.3. Multi-headed attention

Each attention block might contain more than one parallel attention layer, aka attention "head". In theory this allows for different layers to learn various kinds of relationships in the

input and is a very successful tool e.g. in natural language processing [5].

The outputs of all h attention heads are concatenated and multiplied by a weight matrix $W^O \in \mathbb{R}^{hd_v \times d_{model}}$ to squeeze them back into $d_I \times d_{model}$ shape before passing them to the next layer.

To offset the additional computational cost related with multiple attention heads, we set d_k and d_v to d_{model}/h.

1.4. Training

1.4.1. Additional input

[pk: what is a covariate?] [mn: I changed the name since it's not really necessary to call these covariates.]

Following [1], we use additional variables as part of the model's input aside the time series itself. There are three kinds of additional input we use:

1. Numerical information about time the input was recorded (year, month, day-of-the-week, hour-of-the-day, minute-of-the-hour).

2. Identifier of the time series.

3. Input's distance from the sequence's start.

We embed all of the additional variables to the same dimension, sum those embeddings and normalize them. Finally, we add the result to our model's (embedded) input and normalize it as well.

1.4.2. Scaling and weighted sampling

The time series in a single dataset can vary in their ranges, which might obstruct model's learning process. Due to network's non-linearities, it needs to size down the input in its first layer and then reverse this scaling on the output ([7]). Therefore, following [1] and [7] we scale the input and output by respectively dividing and multiplying it by a scale factor v_i:

$$v_i = r + \frac{\sum_{t=1}^{T} y_t^{(i)}}{T},$$

where T is $y^{(i)}$'s length and r is a regularization term.

Furthermore, we adapt a weighted sampling technique in our training, similarly to [1] and [7]. Its main purpose is to counteract the model's tendency to underfit data of a large scale, since it occurs more rarely in the datasets. We calculate the weights by normalizing the scale factor v_i.

1.5. Metrics

1.5.1. Quantile loss

Quantile loss is a metric designed to evaluate and/or train models designed with quantile prediction in mind. We define ρ-quantile loss as follows:

$$QL_\rho(\hat{y}, y) = max(\rho(y - \hat{y}), (\rho - 1)(y - \hat{y})),$$

where $\hat{y}$ denotes the model's prediction for the ρ quantile and y is the ground truth value.

This metric is characterized by a stronger penalization of over-predictions for ρ smaller than 0.5, and under-predictions for ρ greater than 0.5. [pk: good intuition]

1.5.2. Risk metric

$Risk_\rho$ is a type of normalized quantile loss metric defined as follows for $\rho \in (0, 1)$:

$$Risk_\rho(\hat{\mathbf{y}}, \mathbf{y}) = 2 \frac{\sum_{i,t} QL_\rho(\hat{y}_t^{(i)}, y_t^{(i)})}{\sum_{i,t} |y_t^{(i)}|}$$

It is very popular in time series prediction papers that use the Transformer (e.g. in [1], [7]), which lead us to choose it as the metric we use in the model's evaluation. [pk: sure, but why use it specifically in our case, over pointwise metrics?] [mn: I added some details here:] Furthermore, quantile loss allows us to assess how well our model estimates series' extreme values. It is useful for applications like electricity consumption risk estimation, where one might ask how likely it is in the given time-frame that the power consumption reaches a certain level.

Chapter 2

Approximating a continuous distribution with the Transformer

While the Transformer was already proved to be a very efficient tool for the time series prediction, we believe it has a lot more potential that could be used without changing its architecture. In the past, works like [1] were centered around the idea of every point in a time series being drawn from some simple distribution, e.g. the Gaussian one. Thus, the model's goal was to predict the parameters of these distributions. In contrast, one could apply a network to directly forecast series values, which seems like a much simpler approach. However, one of the main benefits of modeling distributions is how it opens the door to using the risk loss. This is especially useful when one considers fields where asking how likely is the time series to exceed a certain value is of high importance, e.g. high-risk markets or traffic jam predictions.

This thesis explores whether the Transformer's performance on time series prediction can be improved by training it to predict a categorical distribution for each point in the sequence instead of a Gaussian one. Since continuous distributions can be approximated with categorical ones, this should enable our Transformer to model a broad spectrum of distributions. One should also take into consideration that Transformer is very successful in tasks requiring prediction of thousands categories like language translation ([5]). This could potentially indicate that the previous usages of the Transformer for the time series prediction were in fact underutilizing its abilities.

We further hypothesize that this ability to model a wide range of distributions could be very useful for especially difficult or non-standard datasets that simpler models would struggle with. For an instance, it would be impossible to train a Gaussian distribution-based Transformer how to recognize a multimodal distribution or a heavy-tailed one.

2.1. Serial Predictor

The main contribution of our work is a time series forecasting model named Serial Predictor that uses the Transformer at its core and does not modify its architecture. The only introduced changes are to the input preprocessing, as we use a discretization process on it that we call **serialization**. It enables us to turn the input into a series of discrete tokens from the $[0, 1]$ interval. We hypothesize that serializing and embedding the time series before feeding it to a Transformer can make it easier for the model to train, since it helps to normalize the input.

2.1.1. Serialization

We call **serialization** the reversible process of bucketing each value in a real-valued series, i.e. mapping it to a discrete categorical value from a finite number of categories.

Our serialization algorithm requires two parameters: maximum value we expect in the datasets, minimal (usually equal to 0) and the exact number of categories ($vocabulary\_size$). To serialize a series, we first assign intervals of real values uniformly starting with the minimal value to $vocabulary\_size$ categories. Afterwards we bucket every value in the series to its appropriate interval/category and re-scale it to the $[0, 1]$ interval.

Deserialization is the reverse of this process in which we output the minimal value assigned to a specific bucket.

Since the number of categories we use is finite, serialization is bound to lose the precise value of some data points in a series. With enough buckets used, we hope this drop in precision should be negligible in comparison to our gains in the prediction's reliability.

2.1.2. Architecture's summary

We use the Transformer's architecture as it was described in Section 1.2. Our input consists of serialized and embedded time series. The final decoder layer outputs $vocabulary\_size$ logits, each of them corresponding to a single bucket. During the evaluation process we sample the buckets from a categorical distribution given by the values of output logits. Finally, we deserialize the sampled bucket.

We use a category cross entropy loss to train our model.

Chapter 3

Experiments

In our experiments we decided to compare our Serial Predictor with a time series forecasting model that we call Gaussian Predictor. This model uses the same Transformer body as Serial Predictor, but it does not use the serialization on its input and outputs parameters of a Gaussian distribution. We use logistic loss for training the Gaussian Predictor.

3.1. Datasets

I chose two datasets for presenting our results in this thesis: *electricity* and *traffic*. They are both based on real-world data and are easily accessible through the GluonTS ([13]) package. Furthermore, they are used in many of the papers we base our work on ([1], [7]), which makes them a good reference point.

The *electricity* dataset is the electric consumption data collected from 370 customers recorded with a frequency of 1 hour ([1], [7]). The *traffic* dataset consists of the hourly occupancy rates of 963 lanes in San Francisco ([1], [7]). Since *traffic* displays higher differences between patterns observed in weekends and weekdays, it is considered a more demanding dataset and is useful for finding out how well the model can capture both long-term and short-term dependencies ([1]).

3.2. Differences from previous works

During the late stages of our work we contacted the authors of [1] and received their code. We discovered that there was a significant difference in our evaluation processes. Mainly, the author's predictions of multiple time steps during the evaluation were always ran based on historical data points (not autoregressively). Our suspicion is that this effect was not the authors' intention and might lead to differences in the results, since it makes the task easier for the model.

Additionally, the datasets provided by GluonTS ([13]) are not identical to the ones used in the experiments from [1] or [7]. Even though the data's source is the same, its processing was different[1].

Furthermore, we did not include the sparse attention in our experiments. This difference should not have contributed to the gap between Gaussian and Serial Predictors, though.

---

[1]Source: https://github.com/awslabs/gluon-ts/issues/830.

3.3. Experiments setup

We used the Adam optimizer for all our experiments along with a rsqrt decay learning rate scheduler. Every experiment was run as follows:

- We used $n\_steps = 1000000$ training steps, with 1000 warm-up steps ($warmup\_steps = 1000$) and $batch\_size = 16$.

- The evaluation was run every $eval\_every = 1000$ steps over $n\_eval\_batches = 128$ samples.

- Sampled sub-sequences' length was equal to $series\_length = 768$.

We ran a randomized parameter grid search over the following parameters:

- Convolutional attention's kernel size: $conv\_attention\_kernel\_width \in \{1, 2, 3, 6, 9\}$

- Number of heads in the multi-headed attention: $n\_heads \in \{1, 2, 4, 8\}$

- Dropout rate (same in various places): $dropout \in \{0.1, 0.2, 0.3, 0.4, 0.5, 0.7, 0.9\}$

- Most of the layers outputs' size: $d\_model \in \{64, 128, 256, 512\}$

- Number of decoder layers: $n\_layers \in \{2, 4, 6, 8\}$

- Learning rate's scheduler constant: $multifactor.constant \in \{0.001, 0.003, 0.01, 0.03, 0.1\}$

- Usage of weighted sampling: $weighted\_sampling \in \{True, False\}$

- (Serial Predictor only) Maximum value we can predict: $high \in \{2, 5, 10\}$

- (Serial Predictor only) Number of buckets: $vocab\_size \in \{1024, 2048, 4096\}$

We ran every experiment with a 16GB of RAM limitation on a cluster containing various GPUs: GTX 1080 Ti, Titan X, Titan V, RTX 2080 Ti and RTX A6000. There were respectively 120 and 125 experiments with the Serial Predictor and the Gaussian Predictor. This slight difference is due to some of the experiments not being able to fit in the memory limit.

Evaluations were carried out autoregressively over a prediction length of 24 hours.

3.4. Results

Best performing models' comparison

	Gaussian Predictor		Serial Predictor		Enhancing Locality ([1])	
	$R_{0.5}$	$R_{0.9}$	$R_{0.5}$	$R_{0.9}$	$R_{0.5}$	$R_{0.9}$
electricity	0.127	0.087	**0.069**	**0.041**	0.070	0.044
traffic	0.402	0.224	**0.123**	**0.093**	0.139	0.094

Table 3.1: Results of our best performing models' and the ones cited from [1]. The best results for each dataset and metric are marked in bold.

We have not managed to reproduce the results from [1] with our Gaussian Predictor, which we attribute to many differences in our work that we investigated as I described in the Section

3.2. There is, however, a large gap between the performance of our Gaussian Predictor and the Serial one in every evaluated metric and dataset, with the Serial Predictor scoring almost a third of the $R_{0.5}$ metric over the $traffic$ dataset. Furthermore, the Serial Predictor performs better in some categories or comparably to the model cited from [1] despite the differences that we found out in the late study stages that were unfavorable for our evaluation process.

Serial Predictor's performance analysis

	$vocab\_size = 1024$		$vocab\_size = 2048$		$vocab\_size = 4096$	
	$R_{0.5}$	$R_{0.9}$	$R_{0.5}$	$R_{0.9}$	$R_{0.5}$	$R_{0.9}$
electricity	0.074	0.046	**0.069**	**0.044**	**0.069**	0.047
traffic	0.133	0.106	0.135	0.99	**0.123**	**0.093**

Table 3.2: Comparison of Serial Predictor models with different $vocab\_size$.

Finally, we present how different number of discretization buckets ($vocab\_size$) influences Serial Predictor's performance. For the $traffic$ dataset we used the best performing model's hyperparameters and fixed them, changing only the $vocab\_size$. For the $electricity$ dataset we used the second best performing model's parameters, as the first one could only be trained with our memory available with the lowest $vocab\_size$. In general one should expect that the higher $vocab\_size$ we use, the better performance we can achieve, since it allows us to have more precision (granularity) in our predictions.

Our model works best with $vocab\_size = 4096$ on the $traffic$ dataset (the highest one), but in the $electricity$ dataset its performance is hindered with the highest granularity. This is likely caused by over-fitting of a larger model.

Chapter 4

Conclusion

4.1. Summary

This thesis compares the accuracy of a Transformer model architecture used for the time series prediction in two ways. The first one is based on important past time series works and involves using the Transformer to predict the parameters of a Gaussian distribution for each forecasted time point. The second way is our Serial Predictor that approximates a continuous distribution with a categorical one. The early results of our work are very promising but also suggest a need to search for a good reproducible baseline reference in the future research.

4.2. Further work

While the work presented in this thesis uses its own fair baseline model, the further experiments should focus on reproducing or readjusting the results from [1] and crafting a consistent comparison with the past time series research.

Furthermore, it is important to attempt finding datasets that might benefit even more from our approach in comparison to established Transformer-based time series models. Since the Serial Predictor can theoretically be trained to output almost any distribution, we hope that it might perform better on data that has low determinism qualities like the stock or cryptocurrency markets.

This work also suggests that it might be beneficial for any further research applying the Transformer model to try leveraging its ability to successfully predict many categories. This is especially relevant in fields where previous works simplified a task to one where it is easier to apply less complicated, non-neural network approaches.

Bibliography

[1] Shiyang Li, Xiaoyong Jin, Yao Xuan, Xiyou Zhou, Wenhu Chen, Yu-Xiang Wang and Xifeng Yan. *Enhancing the Locality and Breaking the Memory Bottleneck of Transformer on Time Series Forecasting.* arXiv:1907.00235, 2019.

[2] Sepp Hochreiter and Jürgen Schmidhuber. *Long short-term memory.* Neural computation, 9(8):1735–1780, 1997.

[3] Kyunghyun Cho, Bart Van Merriënboer, Dzmitry Bahdanau, and Yoshua Bengio. *On the properties of neural machine translation: Encoder-decoder approaches.* arXiv preprint arXiv:1409.1259, 2014.

[4] Urvashi Khandelwal, He He, Peng Qi, and Dan Jurafsky. *Sharp nearby, fuzzy far away: How neural language models use context.* arXiv preprint arXiv:1805.04623, 2018.

[5] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. *Attention is all you need.* NIPS 2017.

[6] Peter J Liu, Mohammad Saleh, Etienne Pot, Ben Goodrich, Ryan Sepassi, Lukasz Kaiser, and Noam Shazeer. *Generating wikipedia by summarizing long sequences.* arXiv preprint arXiv:1801.10198, 2018.

[7] Valentin Flunkert, David Salinas, and Jan Gasthaus. *Deepar: Probabilistic forecasting with autoregressive recurrent networks.* arXiv preprint arXiv:1704.04110, 2017.

[8] George EP Box and Gwilym M Jenkins. *Some recent advances in forecasting and control.* Journal of the Royal Statistical Society. Series C (Applied Statistics), 17(2):91–109, 1968.

[9] R. Hyndman, A. B. Koehler, J. K. Ord, and R .D. Snyder. *Forecasting with Exponential Smoothing: The State Space Approach.* Springer Series in Statistics. Springer, 2008. ISBN 9783540719182

[10] Nikolay Laptev, Jason Yosinski, Li Erran Li and Slawek Smyl. *Time-series extreme event forecasting with neural networks at Uber.* International Conference on Machine Learning.

[11] J. L. Ticknor. *A bayesian regularized artificial neural network for stock market forecasting.* Expert Systems with Applications, 40(14):5501–5506, 2013.

[12] Hongjie Chen, Ryan A. Rossi, Kanak Mahadik, Sungchul Kim and Hoda Eldardiry. *Graph Deep Factors for Forecasting.* arXiv:2010.07373.

[13] Alexander Alexandrov, Konstantinos Benidis, Michael Bohlke-Schneider, Valentin Flunkert, Jan Gasthaus, Tim Januschowski, Danielle C. Maddix, Syama Rangapuram, David Salinas, Jasper Schulz, Lorenzo Stella, Ali Caner Türkmen and Yuyang Wang. *GluonTS: Probabilistic and Neural Time Series Modeling in Python.* Journal of Machine Learning Research.